


## Lecture 2

# Logistic Regression


By: Sultanova Nazerke

# Supervised Machine Learning Algorithms

Classification




Regression



# Classification Problems:


- Spam/non-spam
- Cancer/non-cancer
- Object classification




# Example: Image Classification

- Classify handwritten digits with ML models
- Each input is an entire image
- Output is digit in the image

9 6 6 5 4 0 7 4 0 1  
3 1 3 4 7 2 7 1 2 1  
1 7 4 2 3 5 1 2 4 4



## Classification



Ideally, data is linearly separable - hard decision boundary

# ML Classification algorithms:

- Perceptron
- Logistic Regression
- SVM
- Naïve Bayes
- Decision Trees
- kNN
- Random Forests
- And much more..



or



# Binary classification

- Spam/non spam:

$$y \in \{0, 1\}$$

0: “negative class” (e.g. not spam)

1: “positive class” (e.g. spam)

# Logistic Regression

- Classification Algorithm
- Designed for binary classification, but can be extend to multiclass
- Uses odds ratio and logit function



# Odds ratio and Logit function

- Odds ratio =  $p/(1-p)$
- $p$  is probability of positive event
- Logit function is  $\log(p/(1-p))$

# Logistic function - sigmoid

$$S(z) = \frac{1}{1 + e^{-z}}$$

## Note

- $s(z)$  = output between 0 and 1 (probability estimate)
- $z$  = input to the function (your algorithm's prediction e.g.  $mx + b$ )
- $e$  = base of natural log

# Logistic function - sigmoid


$$S(z) = \frac{1}{1 + e^{-z}}$$

## Note

- $s(z)$  = output between 0 and 1 (probability estimate)
- $z$  = input to the function (your algorithm's prediction e.g.  $mx + b$ )
- $e$  = base of natural log


## Our new hypothesis!!!

# Sigmoid function




# How to apply sigmoid function?

$p \geq 0.5, \text{class} = 1$   
 $p < 0.5, \text{class} = 0$



## Decision boundary



$$Z = w_0 + w_1x_1 + w_2x_2$$

Predict "y = 1" if  $-3 + x_1 + x_2 \geq 0$

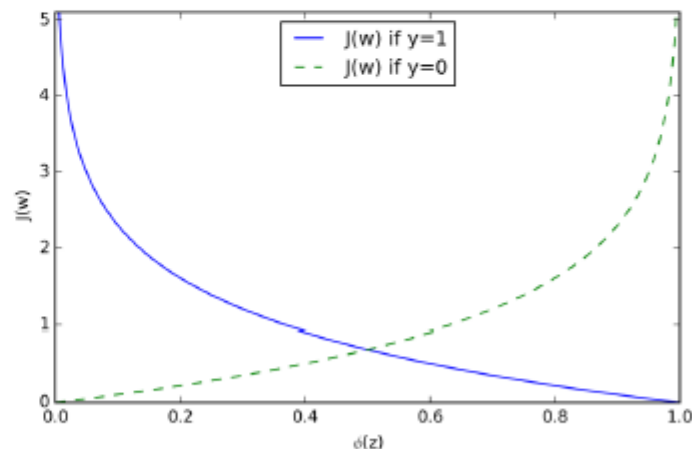
Cost function

Log-likelihood

$$J(w) = \sum_{i=1}^n \left[ -y^i \log(\phi(z^i)) - (1 - y^i) \log(1 - \phi(z^i)) \right]$$

# Learning the weights of the logistic cost function

```
def cost_1(z):  
    return - np.log(sigmoid(z))  
  
def cost_0(z):  
    return - np.log(1 - sigmoid(z))  
  
z = np.arange(-10, 10, 0.1)  
phi_z = sigmoid(z)  
  
c1 = [cost_1(x) for x in z]  
plt.plot(phi_z, c1, label='J(w) if y=1')  
  
c0 = [cost_0(x) for x in z]  
plt.plot(phi_z, c0, linestyle='--', label='J(w) if y=0')  
  
plt.ylim(0.0, 5.1)  
plt.xlim([0, 1])  
plt.xlabel('$\phi(z)$')  
plt.ylabel('J(w)')  
plt.legend(loc='best')  
plt.tight_layout()  
# plt.savefig('./figures/log_cost.png', dpi=300)  
plt.show()
```





# Gradient Descent

$$s'(z) = s(z)(1 - s(z))$$

Which leads to an equally beautiful and convenient cost function derivative:

$$C' = x(s(z) - y)$$

## Note

- $C'$  is the derivative of cost with respect to weights
- $y$  is the actual class label (0 or 1)
- $s(z)$  is your model's prediction
- $x$  is your feature or feature vector.

# Gradient Descent pseudocode

Repeat until convergence {

1. calculate gradient average (just like in linear regression)
2. multiply by learning rate
3. subtract from weights (thetas)

}

# How to know our algorithm works well?

- Test it!
- Calculate accuracy

$\text{accuracy} = (\text{number of correctly labeled}) / (\text{all})$

# Overfitting

- Sometimes model performs well on training data but does not generalize well to unseen data (test data)
- This is overfitting
- If a model suffers from overfitting, the model has a high variance
- This is often caused by a model that's too complex
- Underfitting can also occur (high bias)
- Underfitting is caused by a model's not being complex enough
- Both suffer from low performance on unseen data

# Bias-variance tradeoff





Fig. 1 Graphical illustration of bias and variance.

# How to prevent overfitting??

# Regularization



- Regularization is a way to tune the complexity of the model
- Regularization helps to filter out noise from training data
- As a result, regularization prevents overfitting

# Regularization

- Regularization is a very useful method to handle collinearity (high correlation among features), filter out noise from data, and eventually prevent overfitting.
- The most common form of regularization is the so-called L2 regularization:

$$\frac{\lambda}{2} \|\mathbf{w}\|^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$



# Regularization

- Where  $\lambda$  is the so-called regularization parameter. To apply regularization, we add the regularization term to the cost function, which shrinks the weights:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^i \log(\phi(z^i)) - (1 - \phi(z^i)) \right] + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

# Regularization parameter

- We control how well we fit the training data via the regularization parameter  $\lambda$
- By increasing  $\lambda$ , we increase the strength of regularization
- Sometimes (e.g in scikit-learn), SVM terminology is used

$$C = \frac{1}{\lambda}$$

- i.e. we rewrite the regularized cost function of logistic regression:

$$C \left[ \sum_{i=1}^n \left( -y^i \log(\phi(z^i) - (1 - y^i)) \right) \log(1 - \phi(z^i)) \right] + \frac{\lambda}{2} \|\mathbf{w}\|^2$$


# Visualization of regularization

- Decreasing the value of  $C$  means increasing the regularization strength
- Can be visualized by plotting L2 regularization path for two weights
- Display weights across multiple  $C$  values
- As you see, weights shrink to zero as  $C$  decreased

# Visualization of regularization

```
weights, params = [], []
for c in np.arange(-5., 5.):
    lr = LogisticRegression(C=10.**c, random_state=0)
    lr.fit(X_train_std, y_train)
    weights.append(lr.coef_[1])
    params.append(10**c)

weights = np.array(weights)
plt.plot(params, weights[:, 0],
         label='petal length')
plt.plot(params, weights[:, 1], linestyle='--',
         label='petal width')
plt.ylabel('weight coefficient')
plt.xlabel('C')
plt.legend(loc='upper left')
plt.xscale('log')
# plt.savefig('./figures/regression_path.png', dpi=300)
plt.show()
```




# Assignment 3

## Example: Image Classification

- Classify handwritten digits with ML models
- Each input is an entire image
- Output is digit in the image

9 6 6 5 4 0 7 4 0 1  
3 1 3 4 7 2 7 1 2 1  
1 7 4 2 3 5 1 2 4 4



# Assignment Description

For a full grade you need:

1. Read dataset and visualize one randomly
2. Hypothesis function
3. Cost function
4. Gradient Descent
5. One vs All (Multiclass classification)
6. Calculate accuracy

# Notes for Assignment

- Use vectorized forms –np.arrays because here you have 401 features
- Avoid loops
- To read dataset use:

```
from scipy.io import loadmat  
x = loadmat('test.mat')
```

- **Start earlier!!! No late assignments**