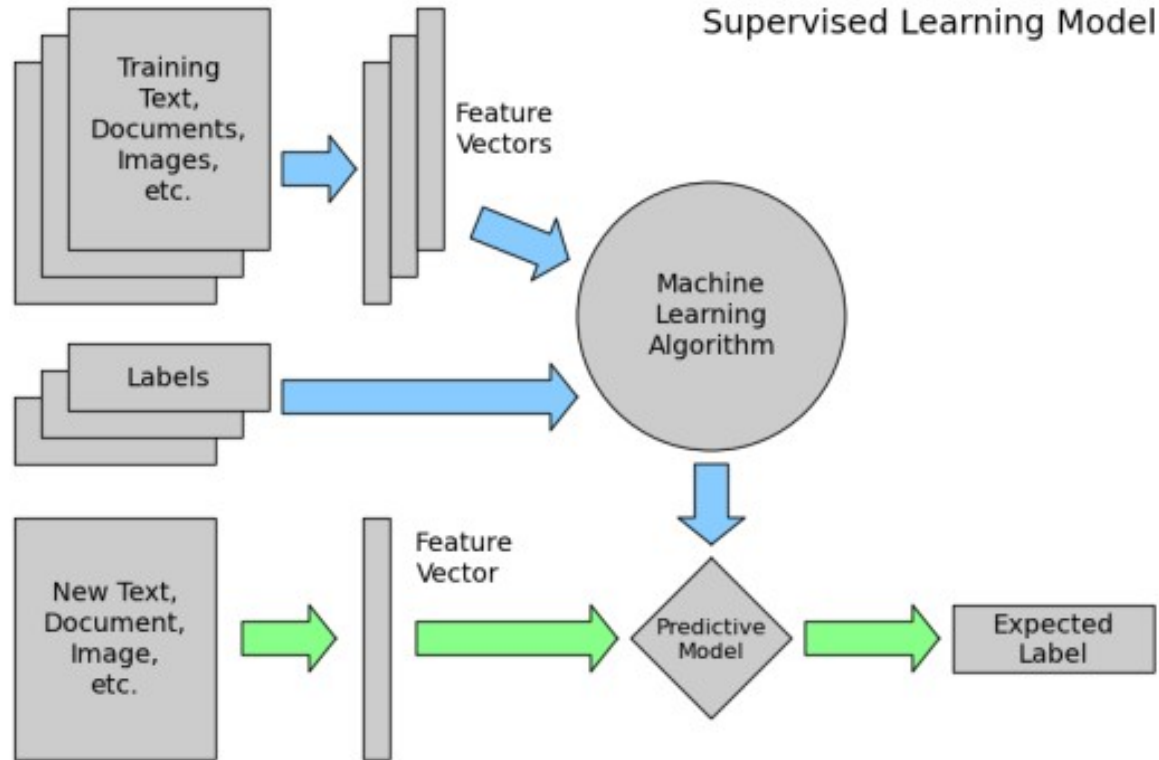


# Lecture 10


## Learning from Text

By Nazerke Sultanova

# Working with text:




# Example:



- How many features needed?

A) 1      B) 2      C) 3      D) 5      E) hard to tell

# Example:



- How many features needed?

A) 1      B) 2      C) 3      D) 5      E) hard to tell

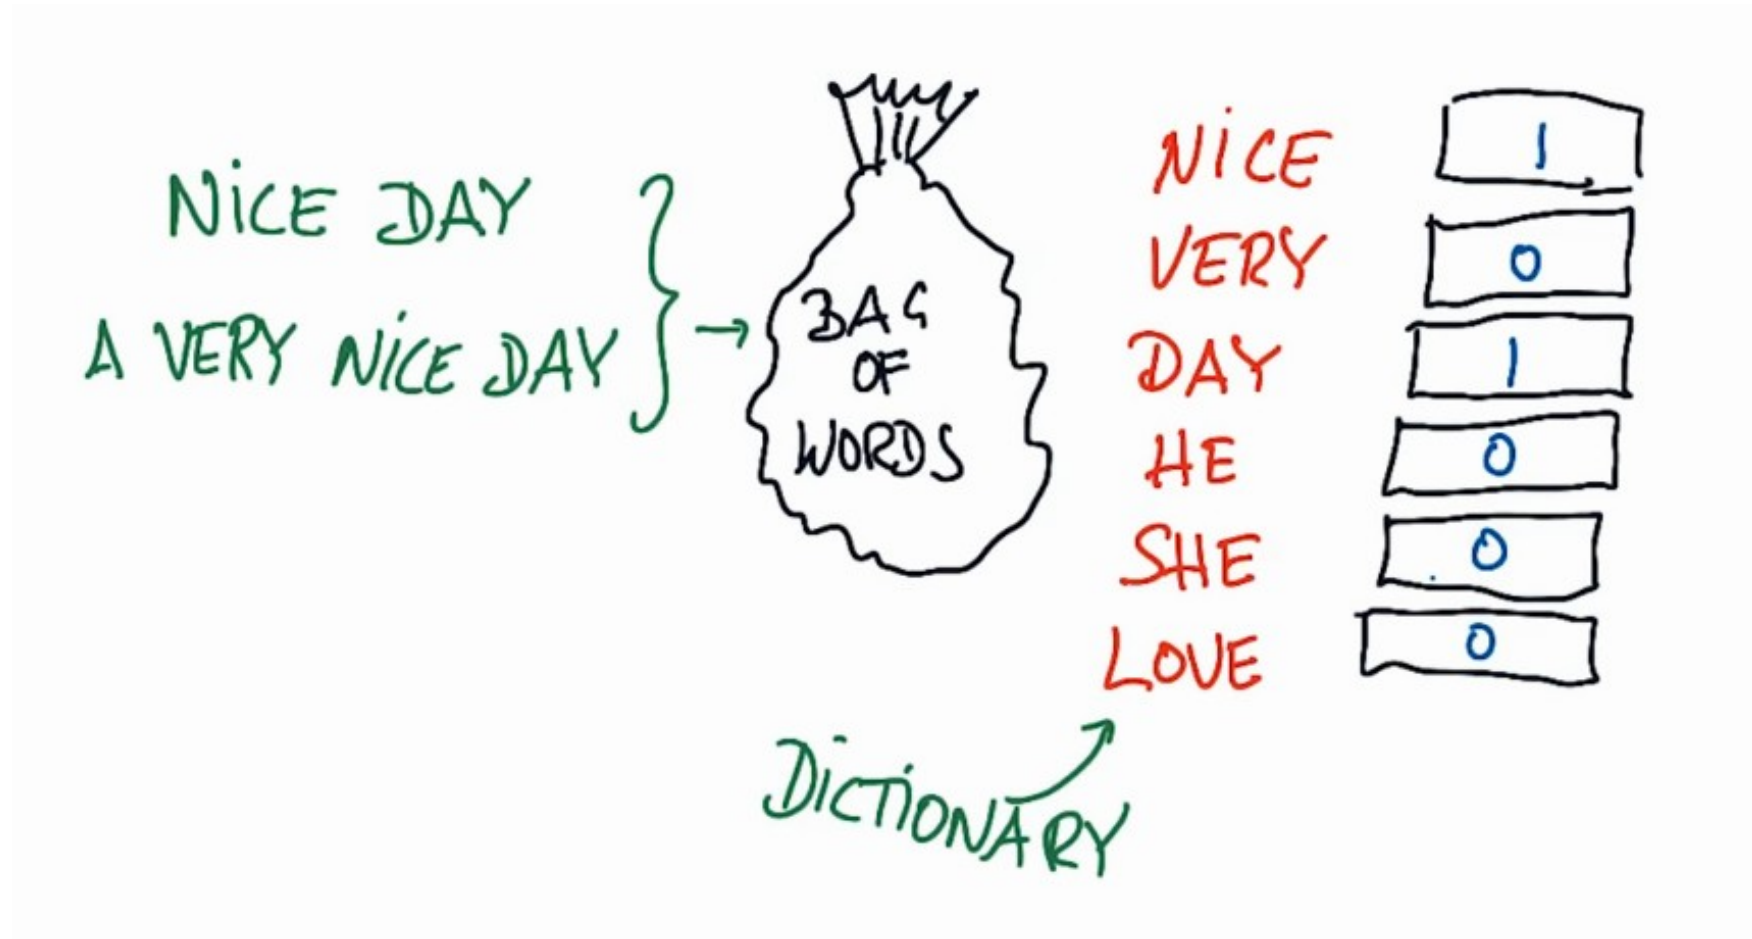
# Learning from text - frequency



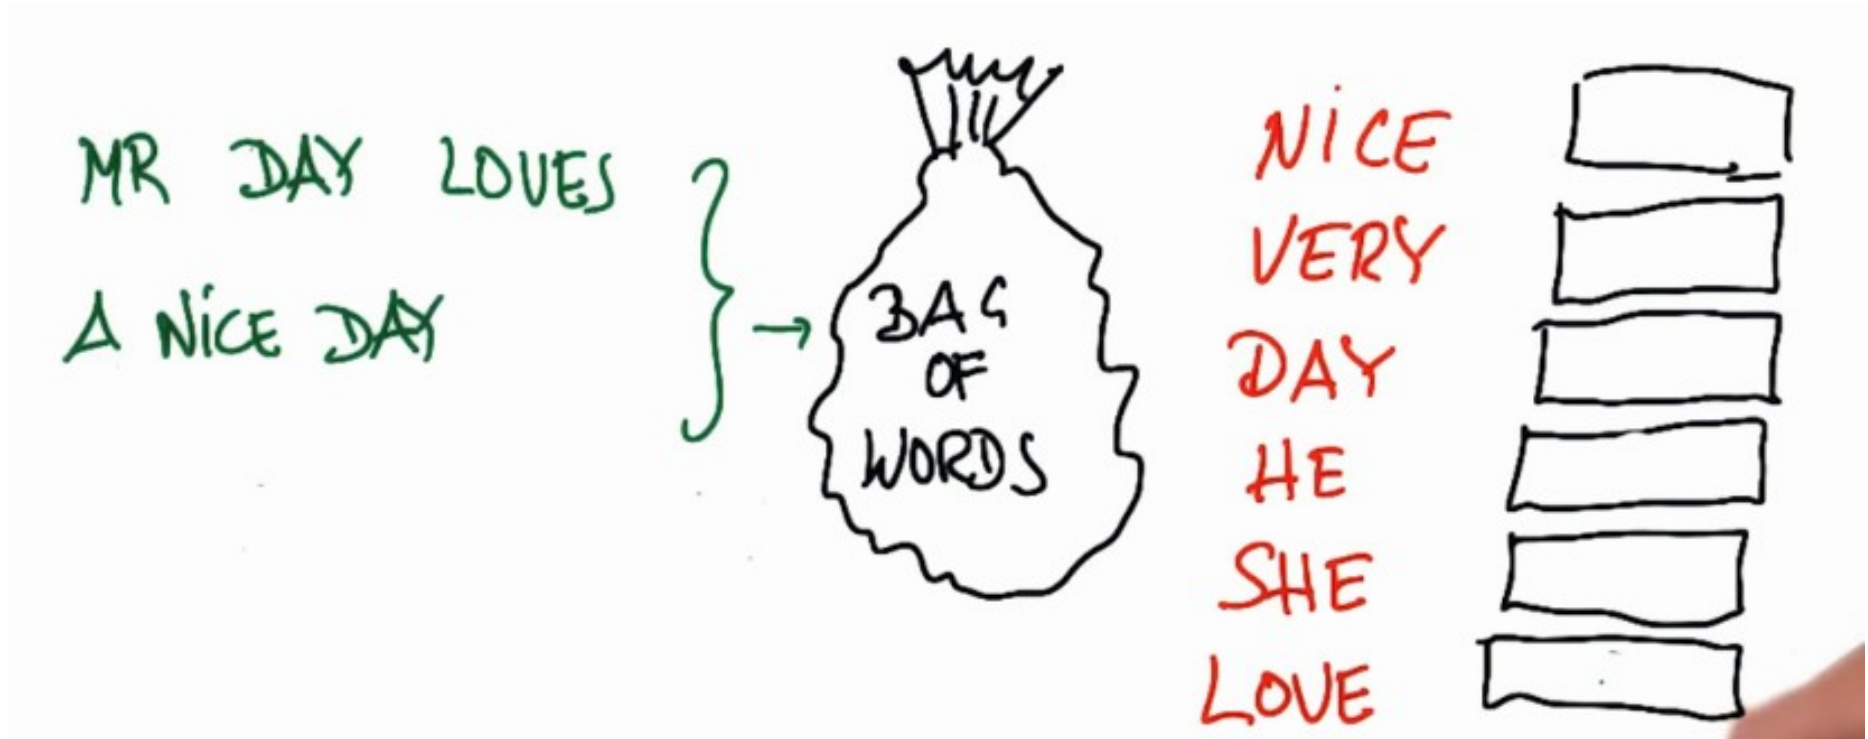
- Fill in frequency vector for a first message:



- Fill in frequency vector for a first message:



- BAG of WORDS works for any text





# BAG OF WORDS properties:

- Does the word order matter?
- Do long phrases give different input vectors?
- Can we handle complex phrases?
  - Ex: “Chicago Bulls”

# BAG OF WORDS properties:

- Does the word order matter? **No**
- Do long phrases give different input vectors?  
**Yes**
- Can we handle complex phrases? **No**
  - Ex: “Chicago Bulls”

- BAG OF WORDS in sklearn - CountVectorizer

```
>>> from sklearn.feature_extraction.text import CountVectorizer
>>> corpus = [
...     'This is the first document.',
...     'This document is the second document.',
...     'And this is the third one.',
...     'Is this the first document?',
... ]
>>> vectorizer = CountVectorizer()
>>> X = vectorizer.fit_transform(corpus)
>>> print(vectorizer.get_feature_names())
['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this']
>>> print(X.toarray())
[[0 1 1 1 0 0 1 0 1]
 [0 2 0 1 0 1 1 0 1]
 [1 0 0 1 1 0 1 1 1]
 [0 1 1 1 0 0 1 0 1]]
```

# Vocabulary: not all words are equal

- Some words contain more information than others
- Find the low information words:  
[‘hi’, ‘Katie’, ‘will’, ‘the’, ‘driving’, ‘machine’, ‘learning’, ‘great’]

# Vocabulary: not all words are equal

- Some words contain more information than others
- Find the low information words:  
[‘hi’, ‘Katie’, ‘will’, ‘the’, ‘driving’, ‘machine’, ‘learning’, ‘great’]

# Vocabulary: not all words are equal

- Some words contain more information than others
- **Stopwords** – low information words that occur highly frequently
- Examples: and, the, I, have, you..
-

Remove stopwords before doing any analysis  
with text!

# Example:

- Stopwords: [the, in, for, you, will, have, be]
- How many words will be removed when we remove stopwords from:  
    ‘Hi Katie, the machine learning class will be great’



# Example:

- Stopwords: [the, in, for, you, will, have, be]
- How many words will be removed when we remove stopwords from:  
    ‘Hi Katie, the machine learning class will be great’
- Answer: 3

# How to get a list of stopwords?

- Python package NLTK
- Natural Language Toolkit (NLTK) is a leading platform for building Python programs to work with human language data

# Importing nltk

```
In [13]: import nltk
          from nltk.corpus import stopwords
          nltk.download()
          sw = stopwords.words('english')
          sw
```

```
Out[13]: ['i',
          'me',
          'my',
          'myself',
          'we',
          'our',
          'ours',
          'ourselves',
          'you',
          "..." ]
```

# Vocabulary – different words with one root

- Unresponsive
- Response
- Responsivity
- Responsiveness
- Respond

# Vocabulary – different words with one root



# Example of stemming - nltk

```
In [4]: from nltk.stem.snowball import SnowballStemmer  
stemmer = SnowballStemmer("english")  
stemmer.stem('responsive')
```

Out[4]: 'respons'

```
In [5]: stemmer.stem('unresponsive')
```

Out[5]: 'unrespons'

---

```
In [7]: stemmer.stem('responsiveness')
```

Out[7]: 'respons'

# Order of operations in text processing

- Which one do we need to do first?
  - Bag of words
  - Stemming
- Why?

# Order of operations in text processing

- Which one do we need to do first?
  - Bag of words
  - Stemming



# TfIdf Representation

- Tf – term frequency – like bag of words
- Idf – inverse document frequency – weighting by how often word occur in a corpus.
  - Rare words weight higher
  - Help to distinguish different messages

## sklearn.feature\_extraction.text.TfidfVectorizer ¶

### Examples

```
>>> from sklearn.feature_extraction.text import TfidfVectorizer
>>> corpus = [
...     'This is the first document.',
...     'This document is the second document.',
...     'And this is the third one.',
...     'Is this the first document?',
... ]
>>> vectorizer = TfidfVectorizer()
>>> X = vectorizer.fit_transform(corpus)
>>> print(vectorizer.get_feature_names())
['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this']
>>> print(X.shape)
(4, 9)
```